

Great Ideas in Computer Architecture

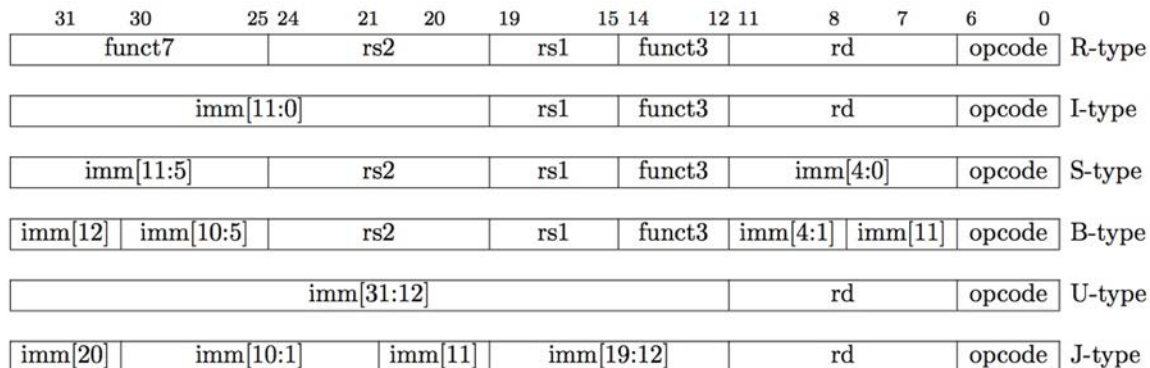
Running a Program: CALL

Instructor: Stephan Kaminsky



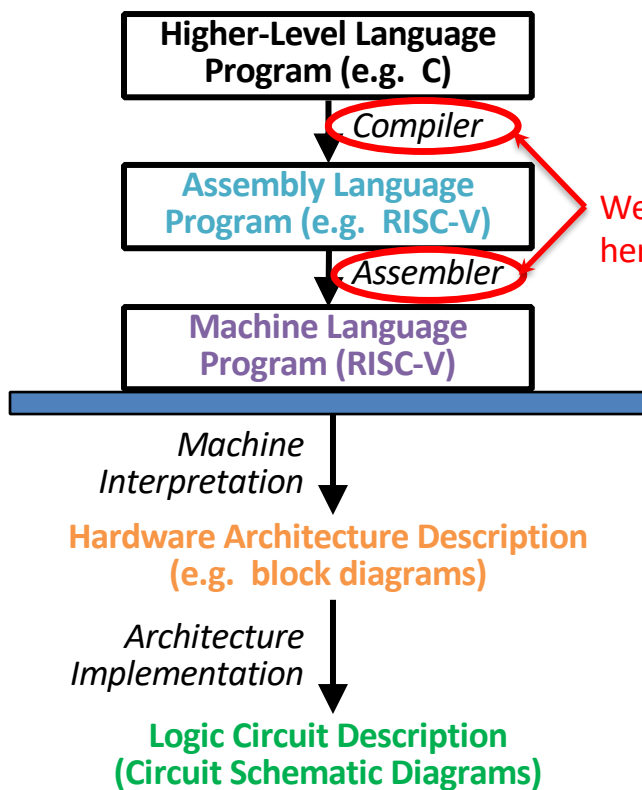
Review

- *Instruction formats* designed to be similar but still capable of handling all instructions



- Branches and Jumps move relative to current address
- Assembly/Disassembly: Use RISC-V Green Sheet to convert

Great Idea #1: Levels of Representation/Interpretation

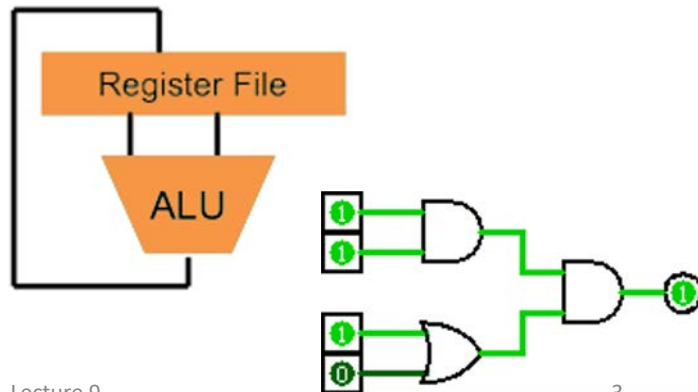


```
temp = v[k];  
v[k] = v[k+1];  
v[k+1] = temp;
```

```
lw    t0, 0(x2)  
lw    t1, 4(x2)  
sw    t1, 0(x2)  
sw    t0, 4(x2)
```

We are
here

```
0000 1001 1100 0110 1010 1111 0101 1000  
1010 1111 0101 1000 0000 1001 1100 0110  
1100 0110 1010 1111 0101 1000 0000 1001  
0101 1000 0000 1001 1100 0110 1010 1111
```



COMPILER

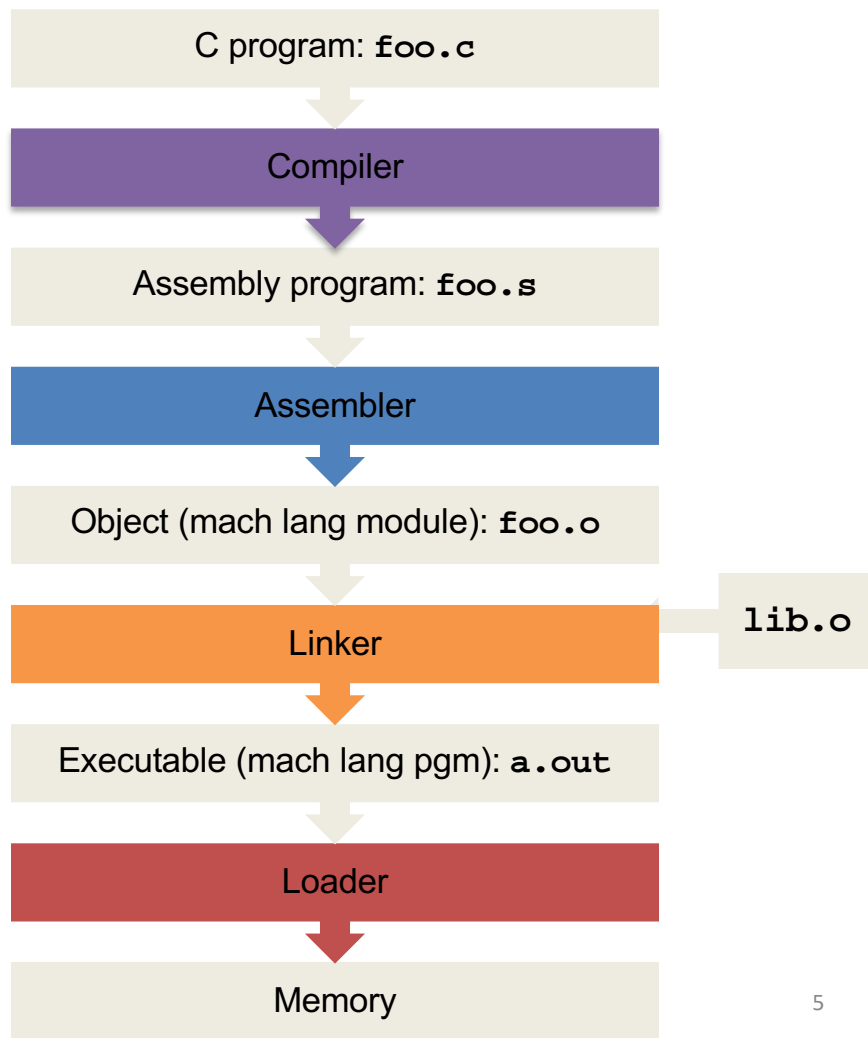
ASSEMBLER

LINKER

LOADER



- Review/Intro
- **Translation**
- Compiler
- Assembler
- Linker
- Loader
- Example



Translation vs. Interpretation (1/3)

- How do we run a program written in a source language?
 - **Interpreter**: Directly executes a program in the source language
 - **Translator**: Converts a program from the source language to an equivalent program in another language
- Directly *interpret* a high level language when efficiency is not critical
- *Translate* to a lower level language when increased performance is desired

Translation vs. Interpretation (2/3)

- Generally easier to write an interpreter
- Interpreter closer to high-level, so can give better error messages (e.g. Python, Venus)
- Interpreter is slower ($\sim 10x$), but code is smaller ($\sim 2x$)
- Interpreter provides instruction set independence: can run on any machine

Translation vs. Interpretation (3/3)

- Translated/compiled code almost always more efficient and therefore higher performance
 - Important for many applications, particularly operating systems
- Translation/compilation helps “hide” the program “source” from the users
 - One model for creating value in the marketplace (e.g. Microsoft keeps all their source code secret)
 - Alternative model, “open source”, creates value by publishing the source code and fostering a community of developers

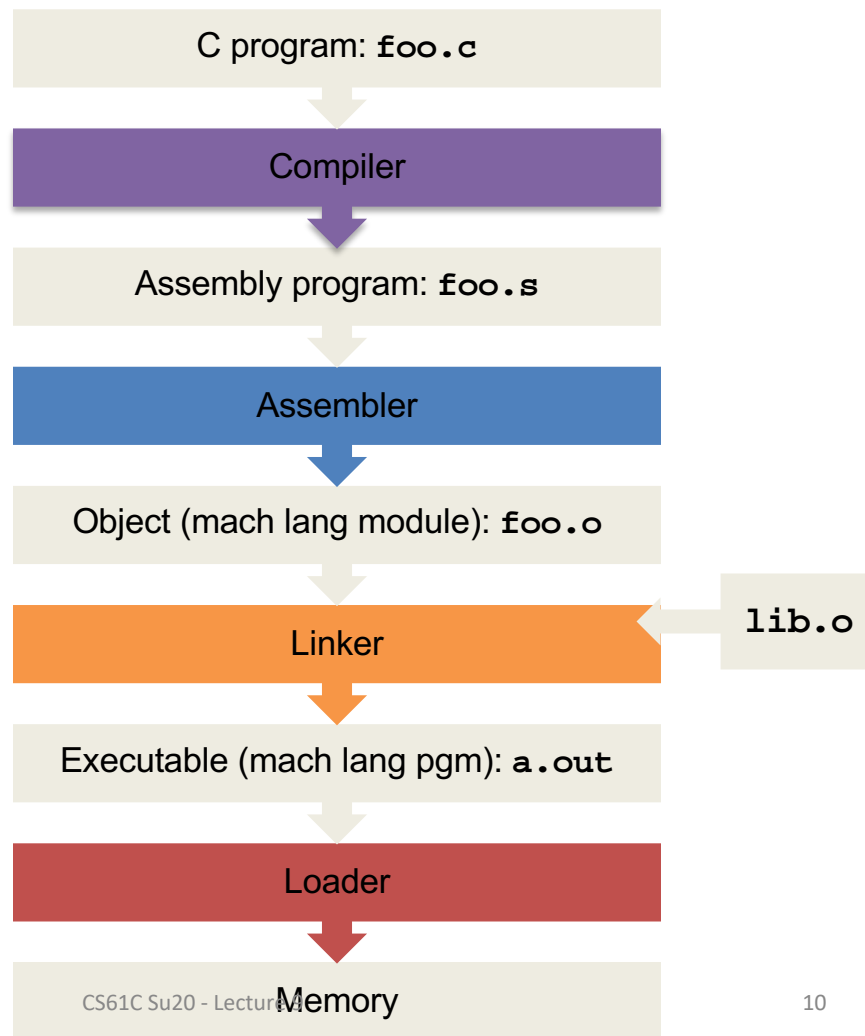
C Translation

- **Recall:** A key feature of C is that it allows you to compile files *separately*, later combining them into a single executable
- What can be accessed across files?
 - Functions
 - Global variables

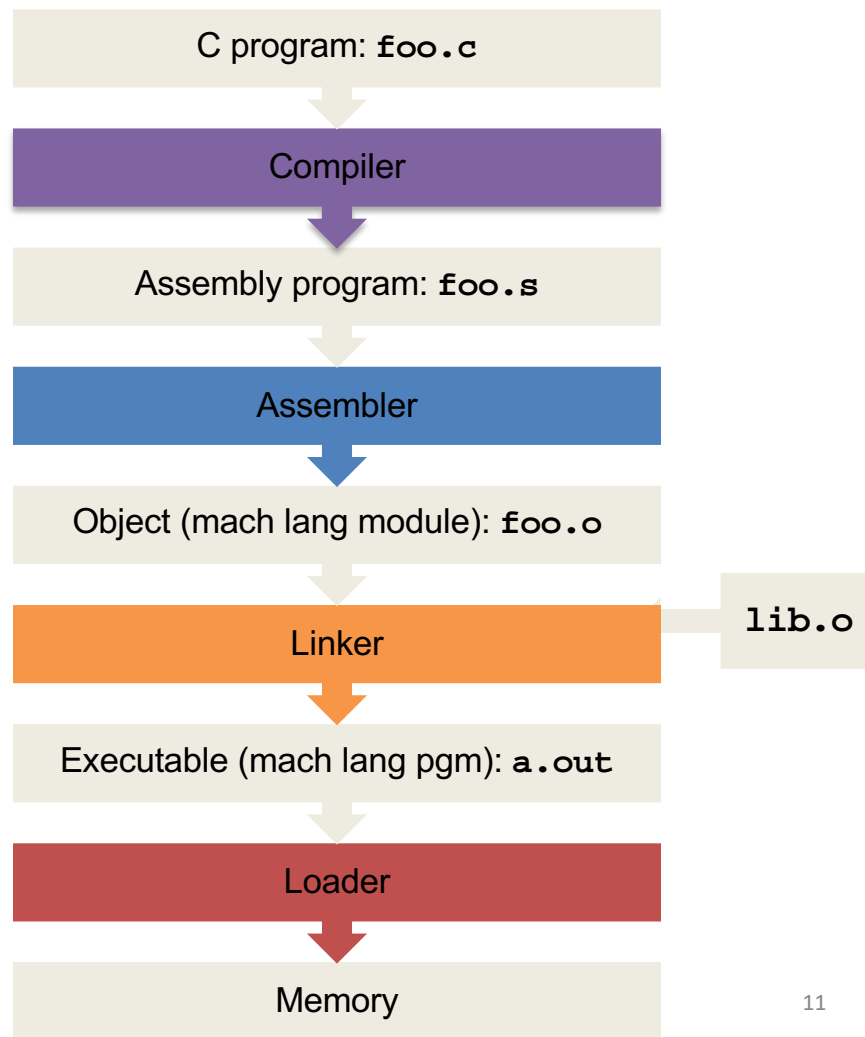
C Translation

Steps to Starting a Program:

- 1) **C**ompiler
- 2) **A**sembler
- 3) **L**inker
- 4) **L**oader



- Review/Intro
- Translation
- **Compiler**
- Assembler
- Linker
- Loader
- Example



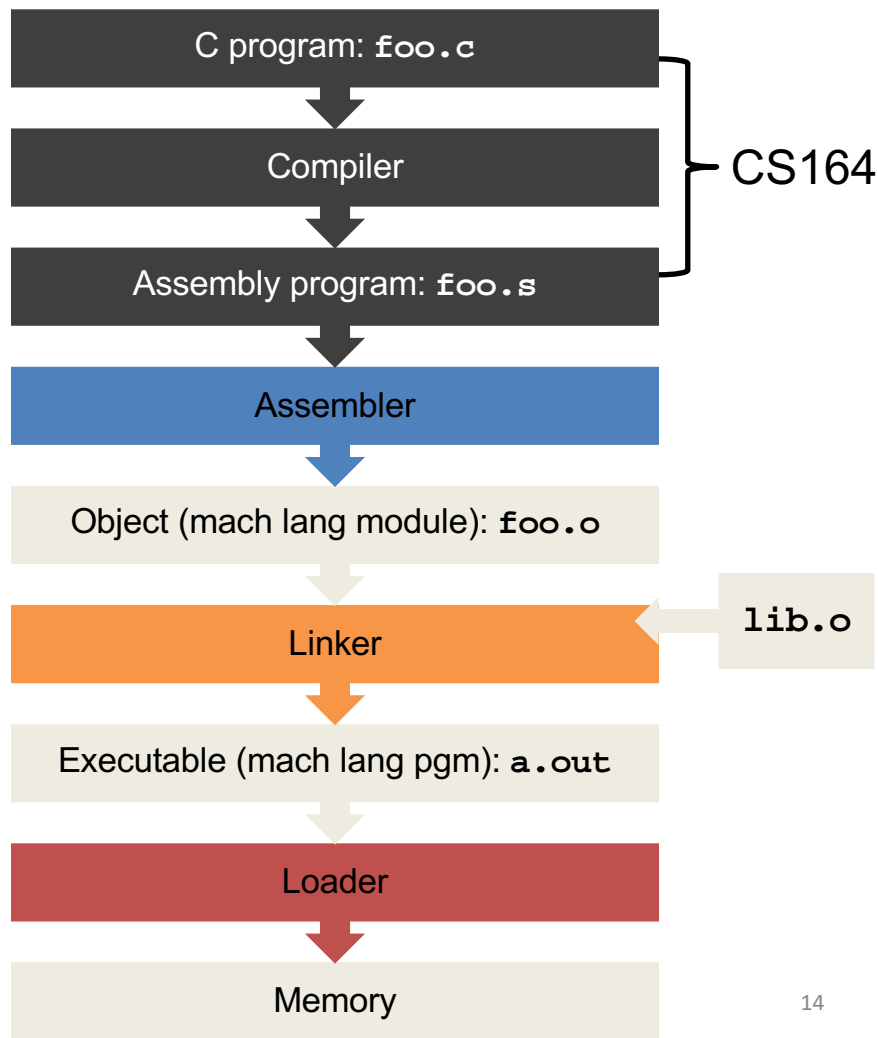
Compiler

- **Input:** Higher-level language (HLL) code (e.g. C, Java in files such as `foo.c`)
- **Output:** Assembly Language Code (e.g. `foo.s` for RISC-V)
- Note that the output may contain pseudo-instructions
- In reality, there's a preprocessor step before this to handle `#directives` but it's not very exciting

Compilers Are Non-Trivial

- There's a whole course about them – CS164
 - We won't go into much detail in this course
 - For the very curious and highly motivated:
<http://www.sigbus.info/how-i-wrote-a-self-hosting-c-compiler-in-40-days.html>
- Some examples of the task's complexity:
 - Operator precedence: $2 + 3 * 4$
 - Operator associativity: $a = b = c$;
 - Determining locally whether a program is valid
 - `if (a) { if (b) { ... /*long distance*/ ... } } }` //extra bracket

- Review/Intro
- Translation
- Compiler
- **Assembler**
- Linker
- Loader
- Example



Assembler

- **Input:** Assembly language code
(e.g. `foo.s` for RISC-V)
- **Output:** Object code (True Assembly),
information tables (e.g. `foo.o` for RISC-V)
—Object file
- Reads and uses **directives**
- Replaces pseudo-instructions
- Produces machine language

Assembler Directives

(For more info, see p.B-5 and B-7 in P&H)

- Give directions to assembler, but do not produce machine instructions
 - .text: Subsequent items put in user text segment (machine code)
 - .data: Subsequent items put in user data segment (binary rep of data in source file)
 - .globl sym: declares `sym` global and can be referenced from other files
 - .ascii str: Store the string `str` in memory and null-terminates it
 - .word $w_1 \dots w_n$: Store the n 32-bit quantities in successive memory words

Pseudo-instruction Replacement

Pseudo

mv t0, t1

neg t0, t1

li t0, imm

not t0, t1

beqz t0, loop

la t0, str

Real

addi t0,t1,0

sub t0, zero, t1

addi t0, zero, imm

xori t0, t1, -1

beq t0, zero, loop

lui t0, str[31:12]

addi t0, t0, str[11:0]

OR

auipc t0, str[31:12]

addi t0, t0, str[11:0]

Producing Machine Language (1/3)

- Simple Cases
 - Arithmetic and logical instructions, shifts, etc.
 - All necessary info contained in the instruction
- What about Branches and Jumps?
 - Branches and Jumps require a *relative address*
 - Once pseudo-instructions are replaced by real ones, we know by how many instructions to branch, so no problem

Producing Machine Language (2/3)

- “Forward Reference” problem
 - Branch instructions can refer to labels that are “forward” in the program:

```
        or    s0, x0,  x0
L1:     slt   t0, x0,  a1
        beq   t0, x0,  L2
        addi  a1, a1, -1
        j     L1
L2:     add   t1, a0,  a1
```

- Solution: Make two passes over the program

Two Passes Overview

- Pass 1:
 - Expands pseudo instructions encountered
 - Remember position of labels
 - Take out comments, empty lines, etc
 - Error checking
- Pass 2:
 - Use label positions to generate relative addresses (for branches and jumps)
 - Outputs the object file, a collection of instructions in binary code

Producing Machine Language (3/3)

- What about jumps to external labels?
 - Requiring knowing a final address
 - Forward or not, can't generate machine instruction without knowing the position of instructions in memory
- What about references to data?
 - `la` gets broken up into `lui` and `addi`
 - These will require the full 32-bit address of the data
- These can't be determined yet, so we create two tables...

Symbol Table

- List of “items” that may be used by other files
 - Each* file has its own symbol table
- What are they?
 - Labels**: function calling
 - Data**: anything in the `.data` section;
variables may be accessed across files
- Keeping track of the labels fixes the forward reference problem

Relocation Table

- List of “items” this file will need the address of later (currently undetermined)
- What are they?
 - Any external **label** jumped to: `jal` or `jalr`
 - internal
 - external (including library files)
 - Any piece of **data**
 - such as anything referenced in the `data` section

Object File Format

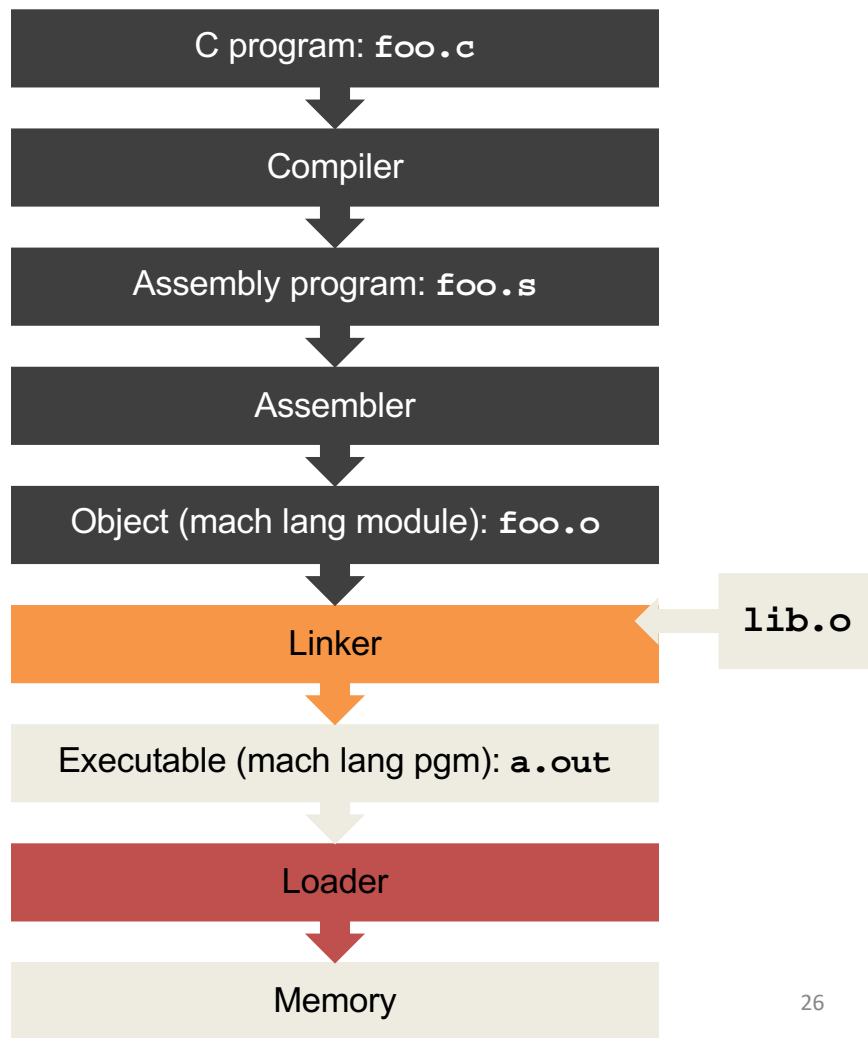
- 1) **object file header**: size and position of the other pieces of the object file
- 2) **text segment**: the machine code
- 3) **data segment**: data in the source file (binary)
- 4) **relocation table**: identifies lines of code that need to be “handled”
- 5) **symbol table**: list of this file’s labels and data that can be referenced
- 6) **debugging information**
 - A standard format is ELF (except MS)

http://www.skyfree.org/linux/references/ELF_Format.pdf

Assembler

- **Input:** Assembly language code
(e.g. `foo.s` for RISC-V)
- **Output:** Object code (True Assembly),
information tables (e.g. `foo.o` for RISC-V)
—Object file
- Reads and uses **directives**
- Replaces pseudo-instructions
- Produces machine language

- Review/Intro
- Translation
- Compiler
- Assembler
- **Linker**
- Loader
- Example




Linker (1/3)

- **Input:** Object Code files, information tables (e.g. `foo.o`, `lib.o` for RISC-V)
- **Output:** Executable Code (e.g. `a.out` for RISC-V)
- Combines several object (`.o`) files into a single executable (“**linking**”)
- **Enables separate compilation of files**
 - Changes to one file do not require recompilation of whole program
 - Old name “Link Editor” from editing the “links” in jump and link instructions

Linker (2/3)

object file 1



object file 2



a.out



Linker (3/3)

- 1) Take text segment from each .o file and put them together
- 2) Take data segment from each .o file, put them together, and concatenate this onto end of text segments
- 3) Resolve References
 - Go through Relocation Table; handle each entry
 - i.e. **fill in all absolute addresses**

Three Types of Addresses

- PC-Relative Addressing (beq, bne, jal)
 - never relocate

External Function Reference (usually jal)

- always relocate

Static Data Reference (often auipc and addi)

- always relocate
- RISC-V often uses auipc rather than lui so that a big block of stuff can be further relocated as long as it is fixed relative to the pc

Absolute Addresses in RISC-V

- Which instructions need relocation editing?
 - J-format: jump/jump and link

XXXXX	jal
-------	-----

- Loads and stores to variables in static area, relative to global pointer

xxx		gp	rd	lw
xx	rs1	gp		sw

- What about conditional branches?

	rs1	rs2	x	beq bne
--	-----	-----	---	------------

- PC-relative addressing preserved even if code moves

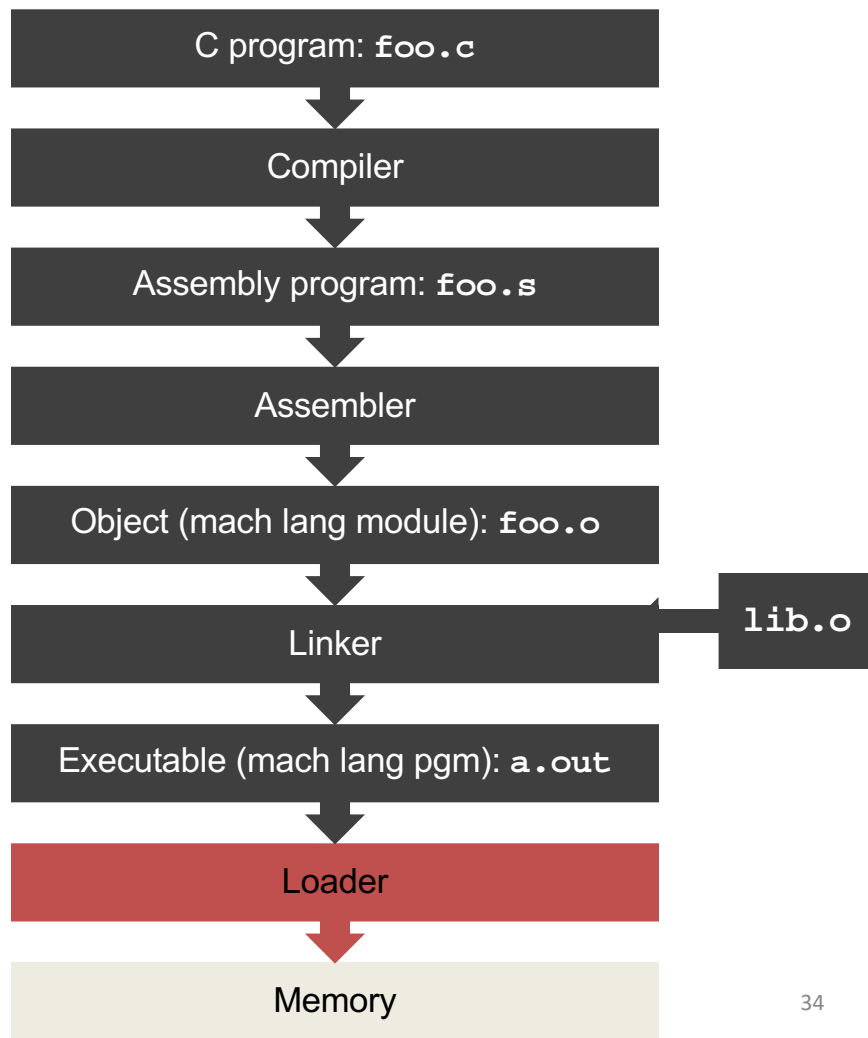
Resolving References (1/2)

- Linker assumes the first word of the first text segment is at **0x10000** for RV32.
 - More later when we study “virtual memory”
- Linker knows:
 - Length of each text and data segment
 - Ordering of text and data segments
- Linker calculates:
 - Absolute address of each label to be jumped to (internal or external) and each piece of data being referenced

Resolving References (2/2)

- To resolve references:
 - 1) Search for reference (data or label) in all “user” symbol tables
 - 2) If not found, search library files (e.g. `printf`)
 - 3) Once absolute address is determined, fill in the machine code appropriately
- Output of linker: executable file containing text and data (plus header)

- Review/Intro
- Translation
- Compiler
- Assembler
- Linker
- **Loader**
- Example



Loader

- **Input:** Executable Code (e.g. `a.out` for RISC-V)
- **Output:** `<program is run>`
- Executable files are stored on disk
- When one is run, loader's job is to load it into memory and start it running
- **In reality, loader is the operating system (OS)**
 - loading is one of the OS tasks

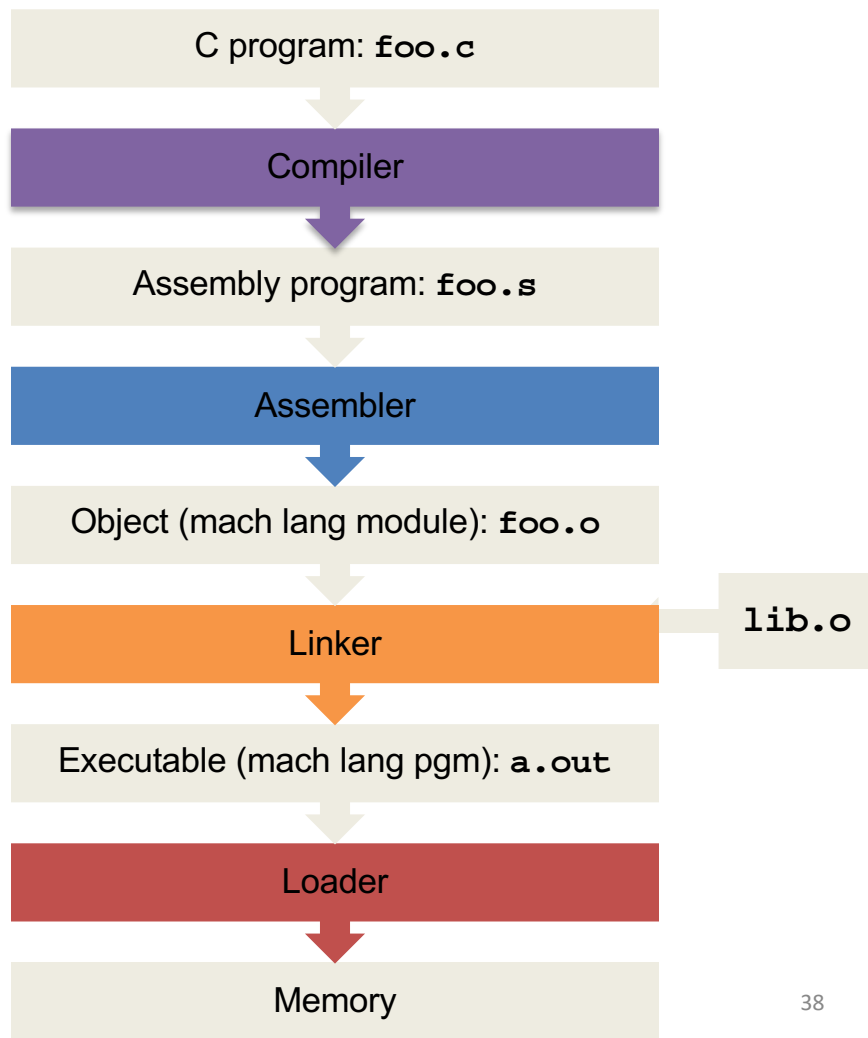
Loader

- 1) Reads executable file's header to determine size of text and data segments
- 2) Creates new address space for program large enough to hold text and data segments, along with a stack segment
<more on this later>
- 3) Copies instructions and data from executable file into the new address space

Loader

- 4) Copies arguments passed to the program onto the stack
- 5) Initializes machine registers
 - Most registers cleared, but stack pointer assigned address of 1st free stack location
- 6) Jumps to start-up routine that copies program's arguments from stack to registers and sets the PC
 - If main routine returns, start-up routine terminates program with the exit system call

- Review/Intro
- Translation
- Compiler
- Assembler
- Linker
- Loader
- **Example**



C.A.L.L. Example

```
#include <stdio.h>  
int main()  
{  
    printf("Hello, %s\n", "world");  
    return 0;  
}
```

Compiled Hello.c: Hello.s

```
.text
    .align 2
    .globl main
main:
    addi sp,sp,-16
    sw   ra,12(sp)
    lui  a0,%hi(string1)
    addi a0,a0,%lo(string1)
    lui  a1,%hi(string2)
    addi a1,a1,%lo(string2)
    call printf
    lw   ra,12(sp)
    addi sp,sp,16
    li   a0,0
    ret
    .section .rodata
    .balign 4
string1:
    .string "Hello, %s!\n"
string2:
    .string "world"
```

```
# Directive: enter text section
# Directive: align code to 2^2 bytes
# Directive: declare global symbol main
# label for start of main
# allocate stack frame
# save return address
# compute address of
#   string1
# compute address of
#   string2
# call function printf
# restore return address
# deallocate stack frame
# load return value 0
# return
# Directive: enter read-only data
section
# Directive: align data section to 4
bytes
# label for first string
# Directive: null-terminated string
# label for second string
# Directive: null-terminated string
```


Assembled Hello.s: Linkable Hello.o

```
00000000 <main>:
0:   ff010113 addi sp,sp,-16
4:   00112623 sw ra,12(sp)
8:   00000537 lui a0,0x0 # addr placeholder
c:   00050513 addi a0,a0,0 # addr placeholder
10:  000005b7 lui a1,0x0 # addr placeholder
14:  00058593 addi a1,a1,0 # addr placeholder
18:  00000097 auipc ra,0x0 # addr placeholder
1c:  000080e7 jalr ra # addr placeholder
20:  00c12083 lw ra,12(sp)
24:  01010113 addi sp,sp,16
28:  00000513 addi a0,a0,0
2c:  00008067 jalr ra
```

Linked Hello.o: a.out

```
000101b0 <main>:
  101b0: ff010113 addi sp,sp,-16
  101b4: 00112623 sw ra,12(sp)
  101b8: 00021537 lui a0,0x21
  101bc: a1050513 addi a0,a0,-1520 # 20a10
<string1>
  101c0: 000215b7 lui a1,0x21
  101c4: a1c58593 addi a1,a1,-1508 # 20a1c
<string2>
  101c8: 288000ef jal ra,10450 # <printf>
  101cc: 00c12083 lw ra,12(sp)
  101d0: 01010113 addi sp,sp,16
  101d4: 00000513 addi a0,0,0
  101d8: 00008067 jalr ra
```

Summary

- **Compiler** converts a single HLL file into a single assembly file `.c → .s`
- **Assembler** removes pseudo-instructions, converts what it can to machine language, and creates a checklist for linker (relocation table) `.s → .o`
 - Resolves addresses by making 2 passes (for internal forward references)
- **Linker** combines several object files and resolves absolute addresses `.o → .out`
 - Enable separate compilation and use of libraries
- **Loader** loads executable into memory and begins execution